

Wooden boatbuilding, Lugged steel bicycle frames, computers. The usual.

Monday, 30 January 2017

Measuring a MOSFET Power Amplifier.



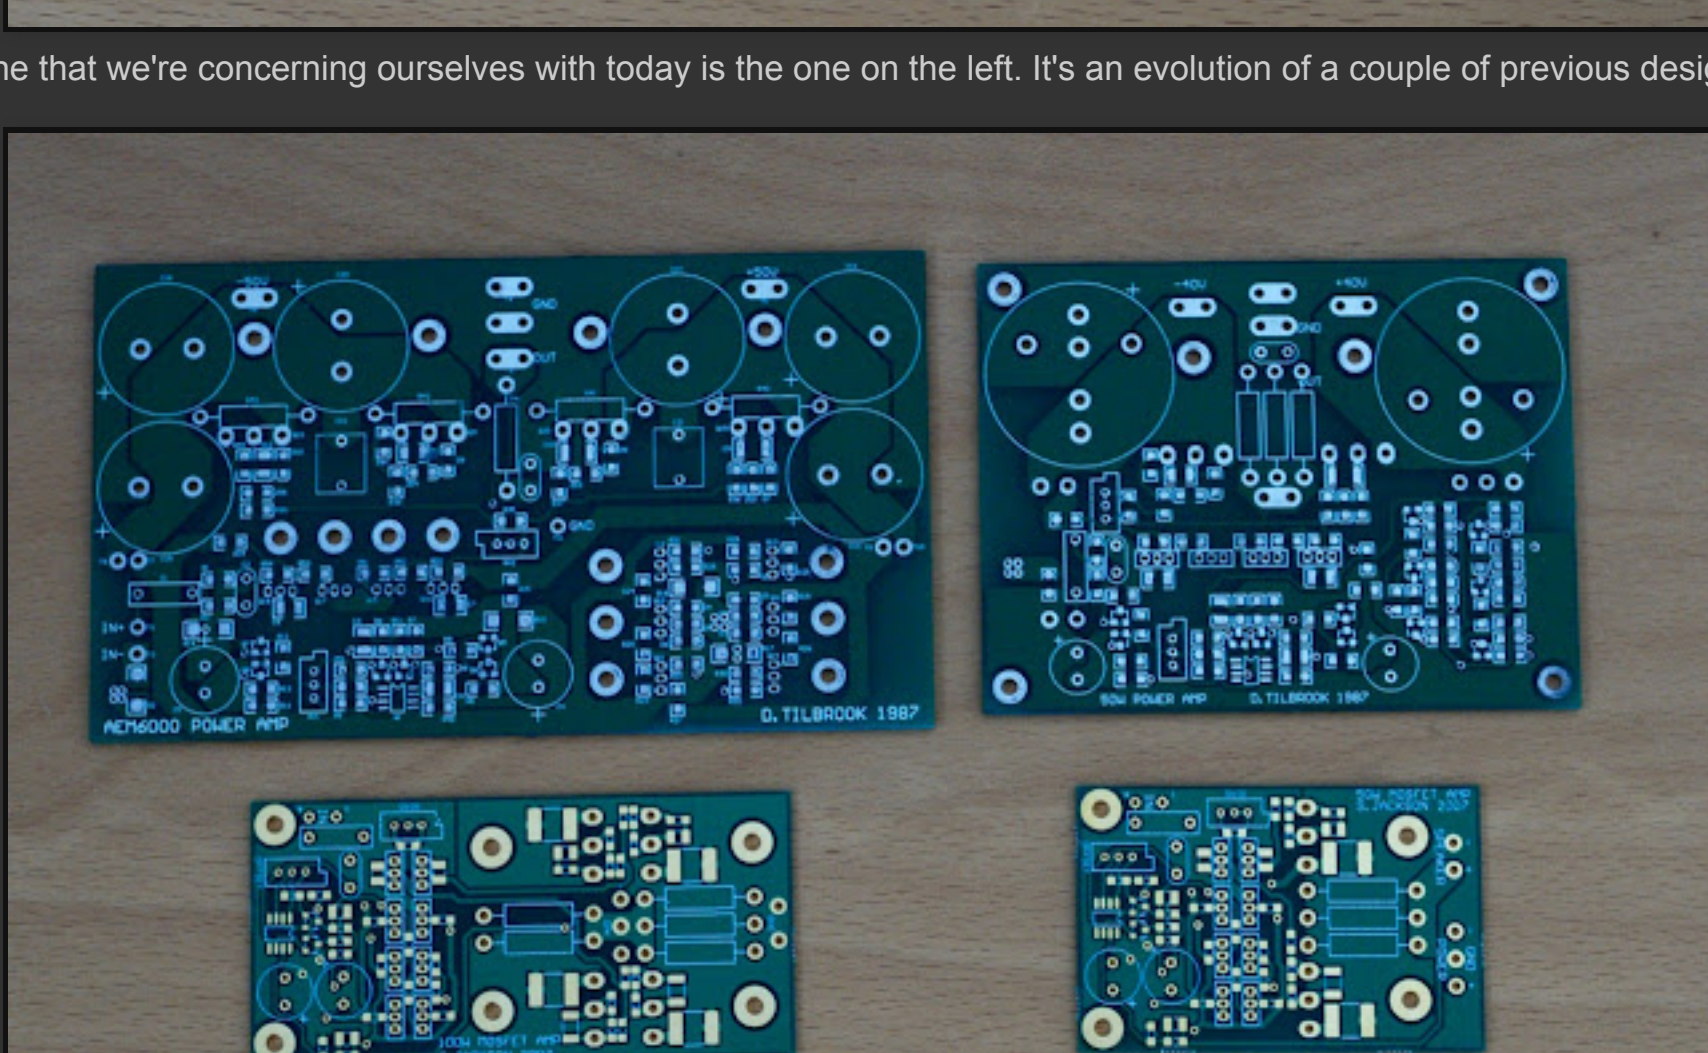
I'm on a bit of an audio kick of late. It's all about building some speakers and an amp for my study, as I stole the speakers and amp from here (some nice home-brew 2-way bookshelves with Vifa P13W1 mid-woofers and P22TC tweeters, and an old NAD 3020B) for my bedroom. Anyway, I've got a pre-existing amp design, based on David Tilbrook's AEM6000, and I've ordered the drivers to build a pair of Lynn Olsen "Ariel" MTM transmission lines. That'll make some pleasant noise, I think.

On to the amps. Back when I originally designed these amps, some 10 years ago, the distortion analyser I had had a noise floor of 0.0013%. I achieved that fairly easily so moved on. I've just bought a PCB for a super-low distortion oscillator I had on the Linear Tech "AN67" circuit, which uses some scary high-gain opamps wrapped around a resonator to generate an "unmeasurably" pure tone.

Now I've got a sniff of a world beyond 0.001%, I'm very keen to see how far I can go. Part of that's developing a whole new completely over-the-top amplifier (see previous posts), part of that is just playing with my existing amps to see what they're really capable of.

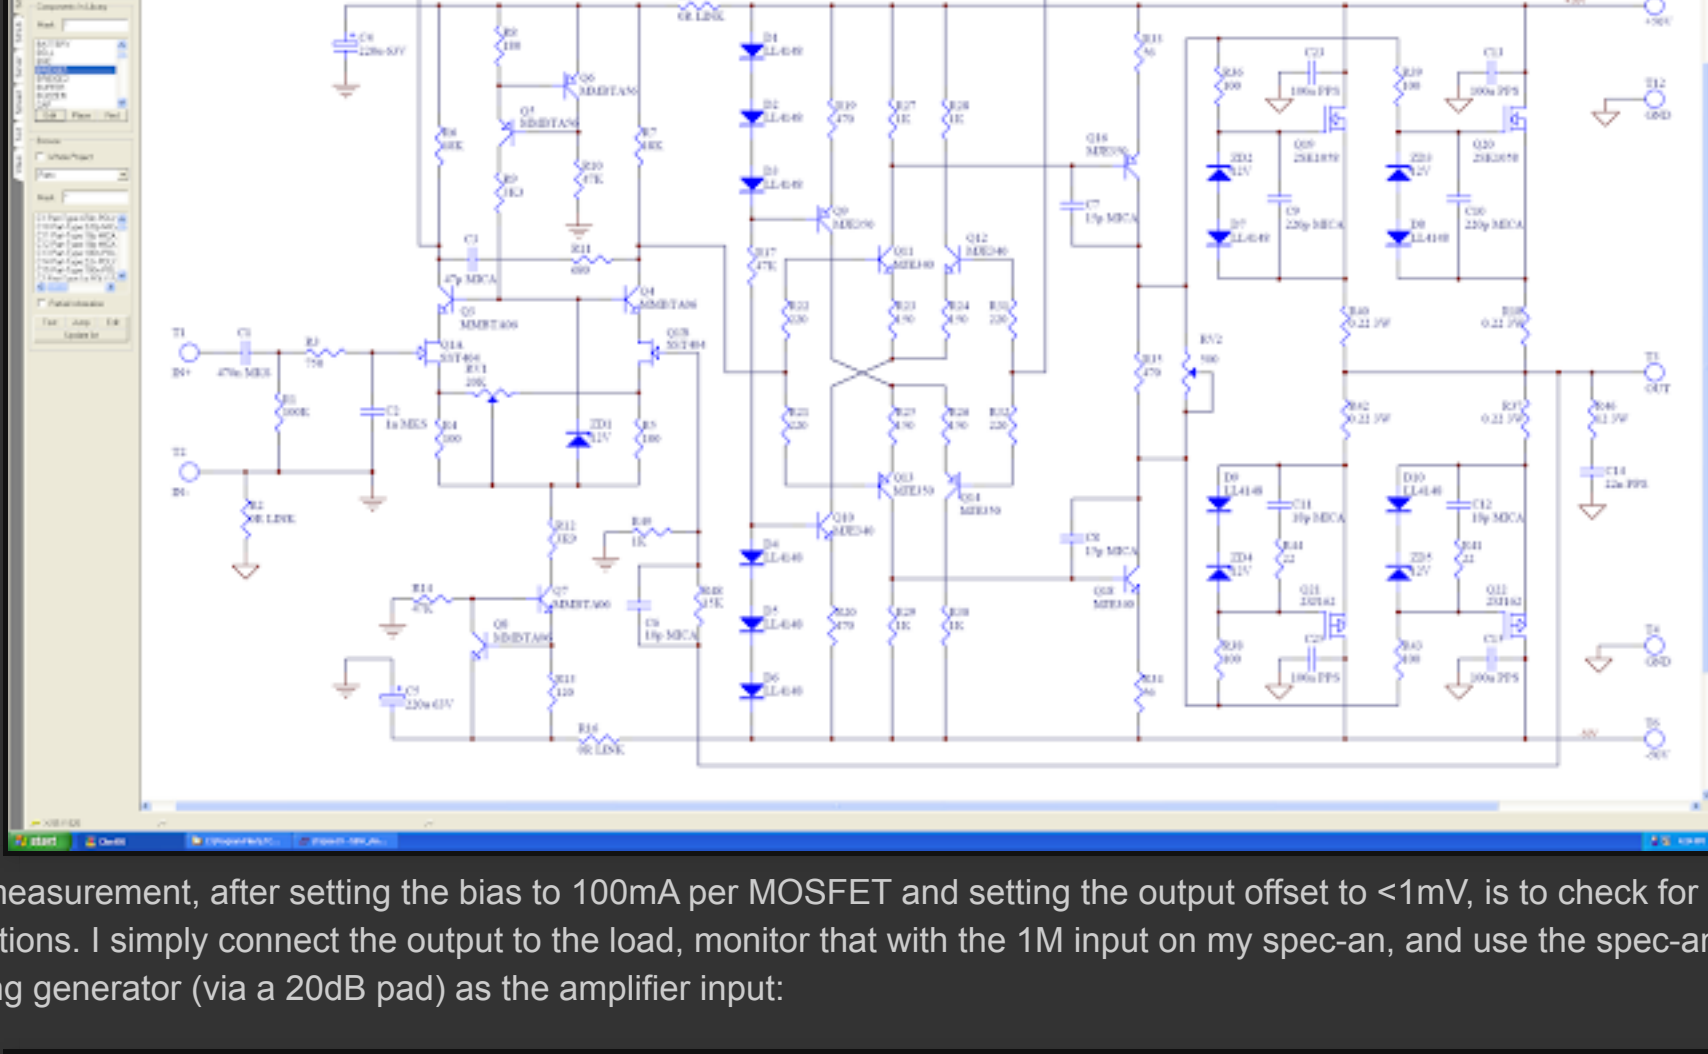
The process is straightforward. Put one together on a heatsink with a dummy load and thrown together power supply using bits from the garage, apply a stimulus from either my mac, my function generator, or my new oscillator, measure what's being put into the load with my CRO and spectrum analyser, tweak, repeat. It's a process that's quite involving.

Here's the amplifier in question. It's a little 50mm x 100mm PCB, implementing a modified version of Tilbrook's AEM6000:

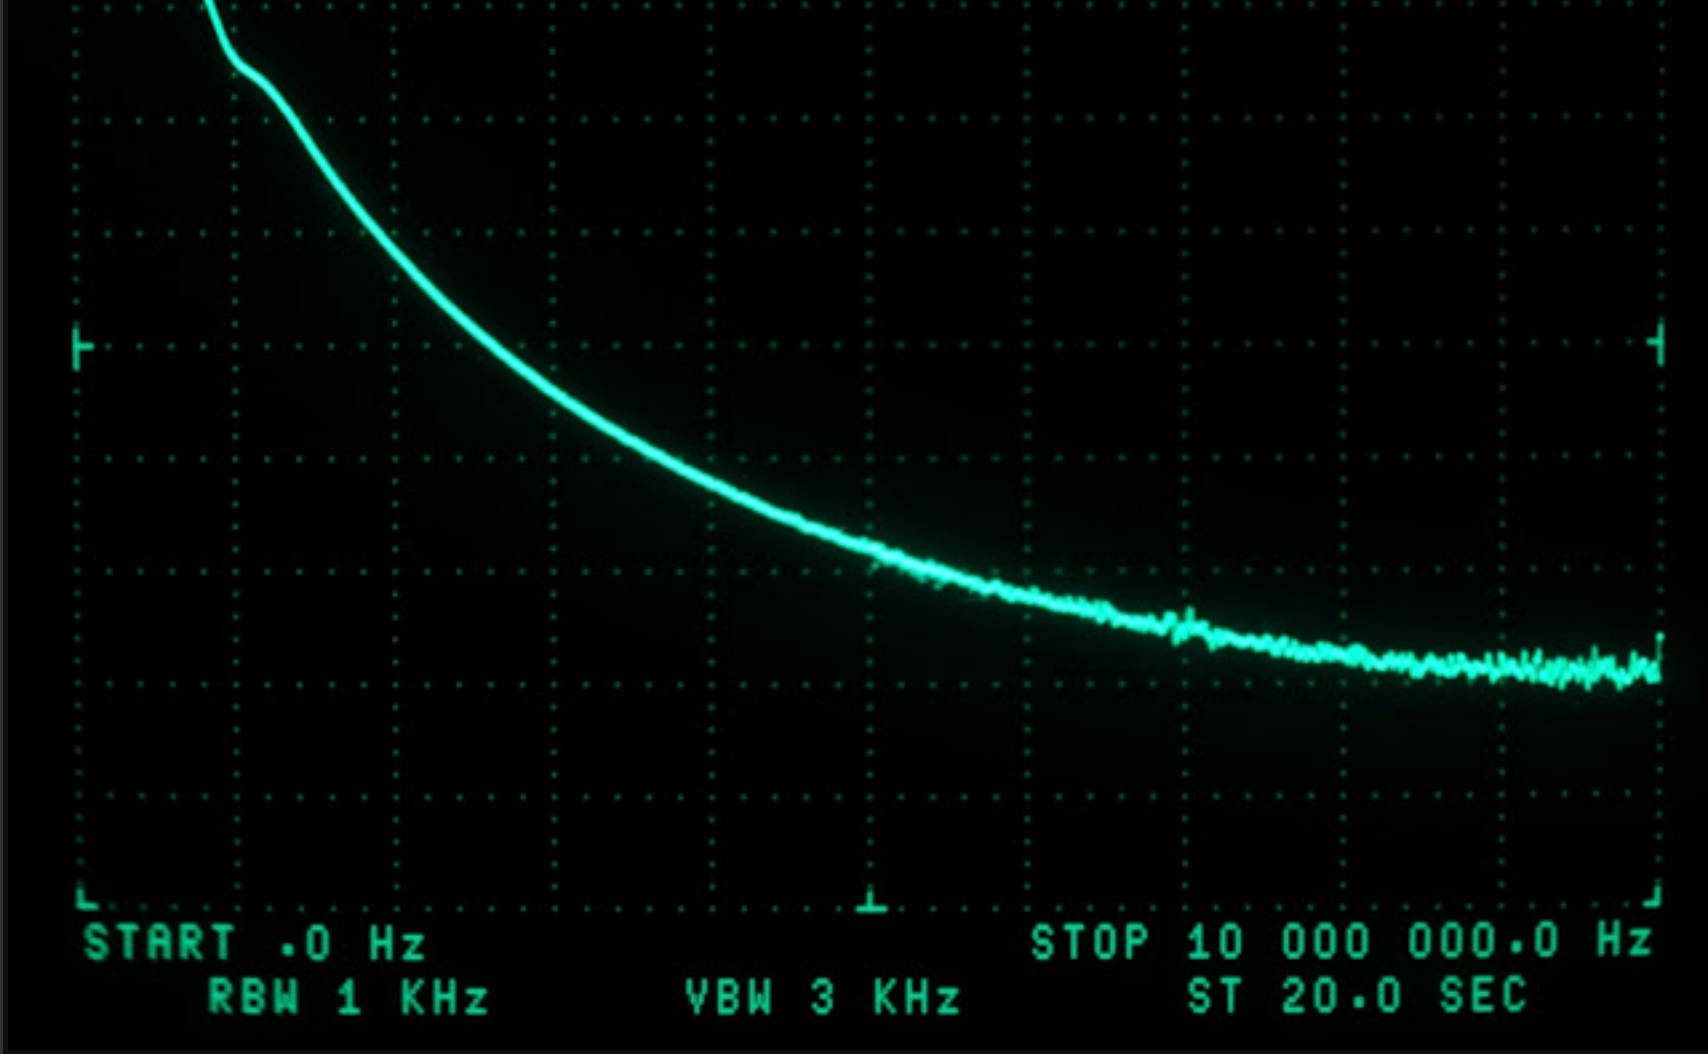


They started with my original "repackaging" of the AEM6000 using flat-pack MOSFETs, shown on the top left. Then I made a "baby" 50W version (top right). While making that I wondered whether I could roll a version that would fit on a 50mm wide heatsink, with off-board filter caps. That resulted in the two smaller boards below. They depart from the original in that the VAS stage is no longer differential-symmetrical, as I simply couldn't see the advantage of the differential structure in Tilbrook's original. Also I made much better use of the back-side of the PCB. Whereas the original has practically nothing on the back bar the MOSFETs, this one has a fairly even spread of parts on each side. That makes it more compact and keeps parasitics down. Despite the small size it's still using 1206 (or MELF 0204) parts.

The schematic for the small 100W flavour is below:



First measurement, after setting the bias to 100mA per MOSFET and setting the output offset to <1mV, is to check for oscillations. I simply connect the output to the load, monitor that with the 1M input on my spec-an, and use the spec-an tracking generator (via a 202B pad) as the amplifier input.




No nasty surprises. Btw lateral MOSFETs are prone to oscillate around 1-2 MHz if the amp isn't well compensated. I find working with LTspice to prototype the compensation is straightforward, and as long as you check the DC conditions on the real thing you'll be good.

While it's connected to the tracking generator, I check the gain flatness across the audio band:



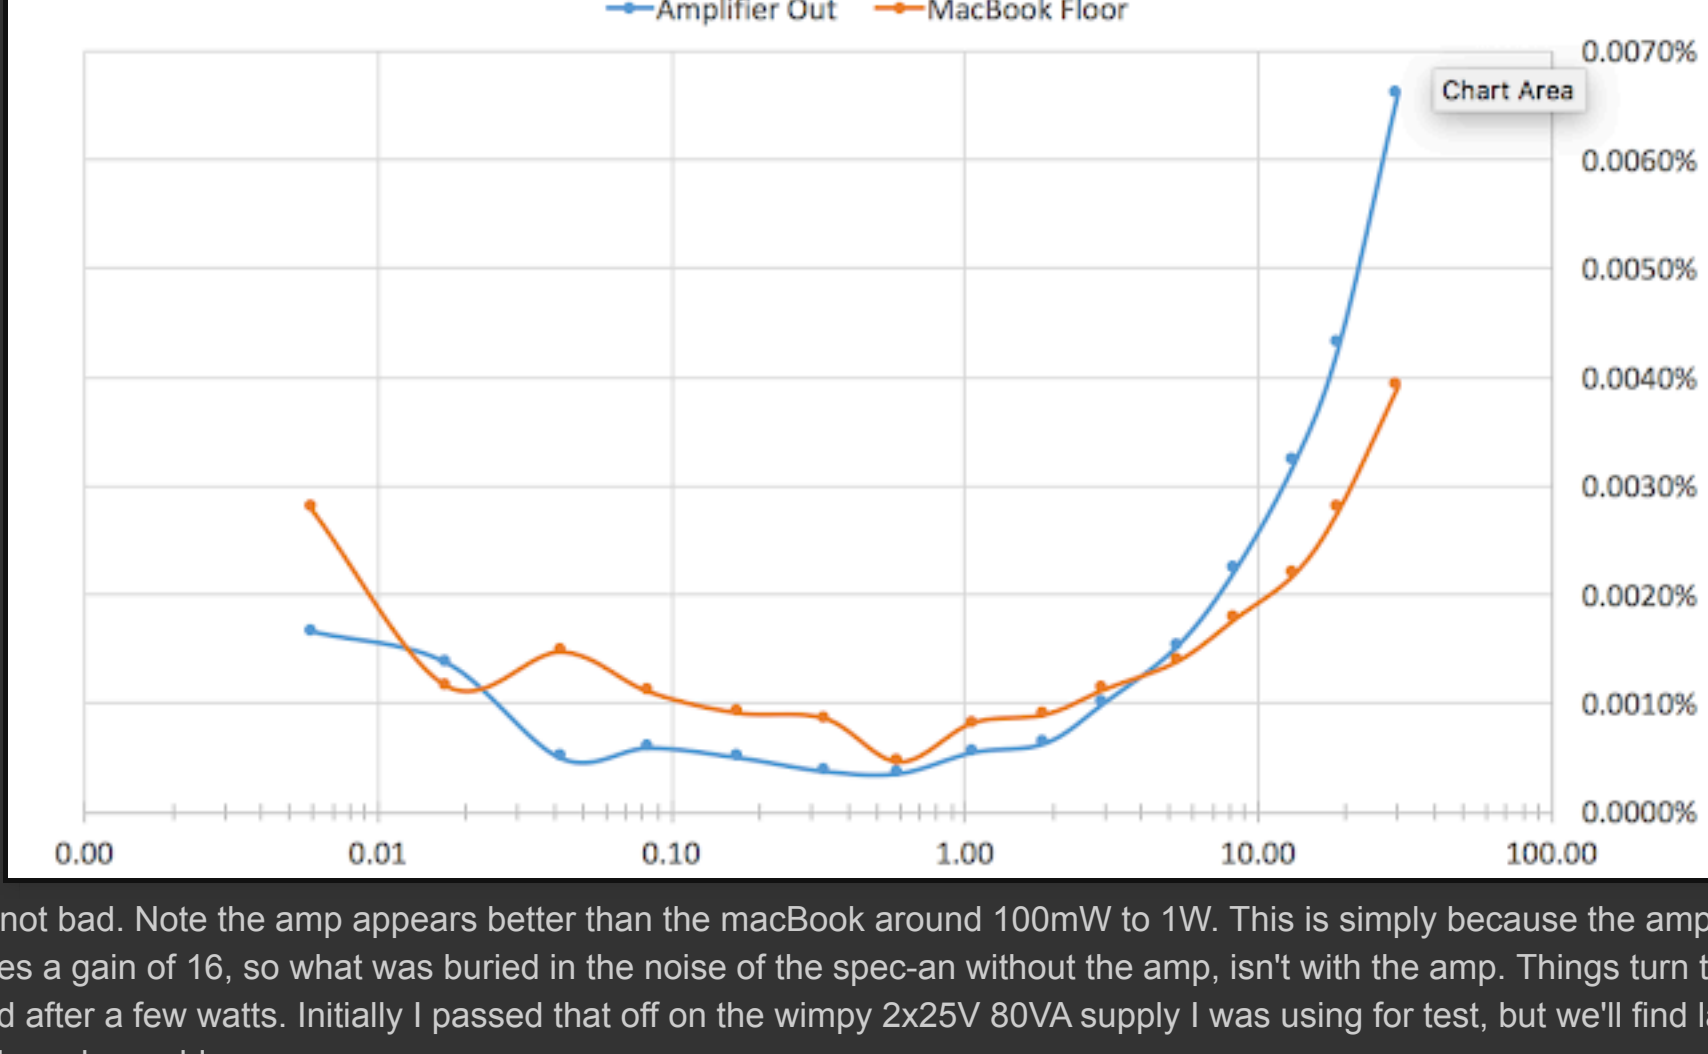
It's hard to tell the spectrum analyser from my amp at the bottom end, as the spectrum analyser has a high-pass at 20Hz. Anyway, it's flat to within 0.1dB or so from there to 20KHz, which should be plenty. Of note is that I measured my NAD3020B, and it's got 6dB of bass boost from 60Hz to about 600Hz, with the bass knob in the middle. It's relatively flat when the bass knob is turned all the way down. I guess that's how you sell amps. That one "sounds better" - ie colours the music.

From here the measurements get increasingly challenging. First the noise floor of the amp. My spectrum analyser has a direct noise readout in nV/rtHz, which is incredibly cool. Pulling the amp input off the tracking generator, terminating it into 50 Ohms, and then winding the spectrum analyser all the way down to find the noise floor, we get this:



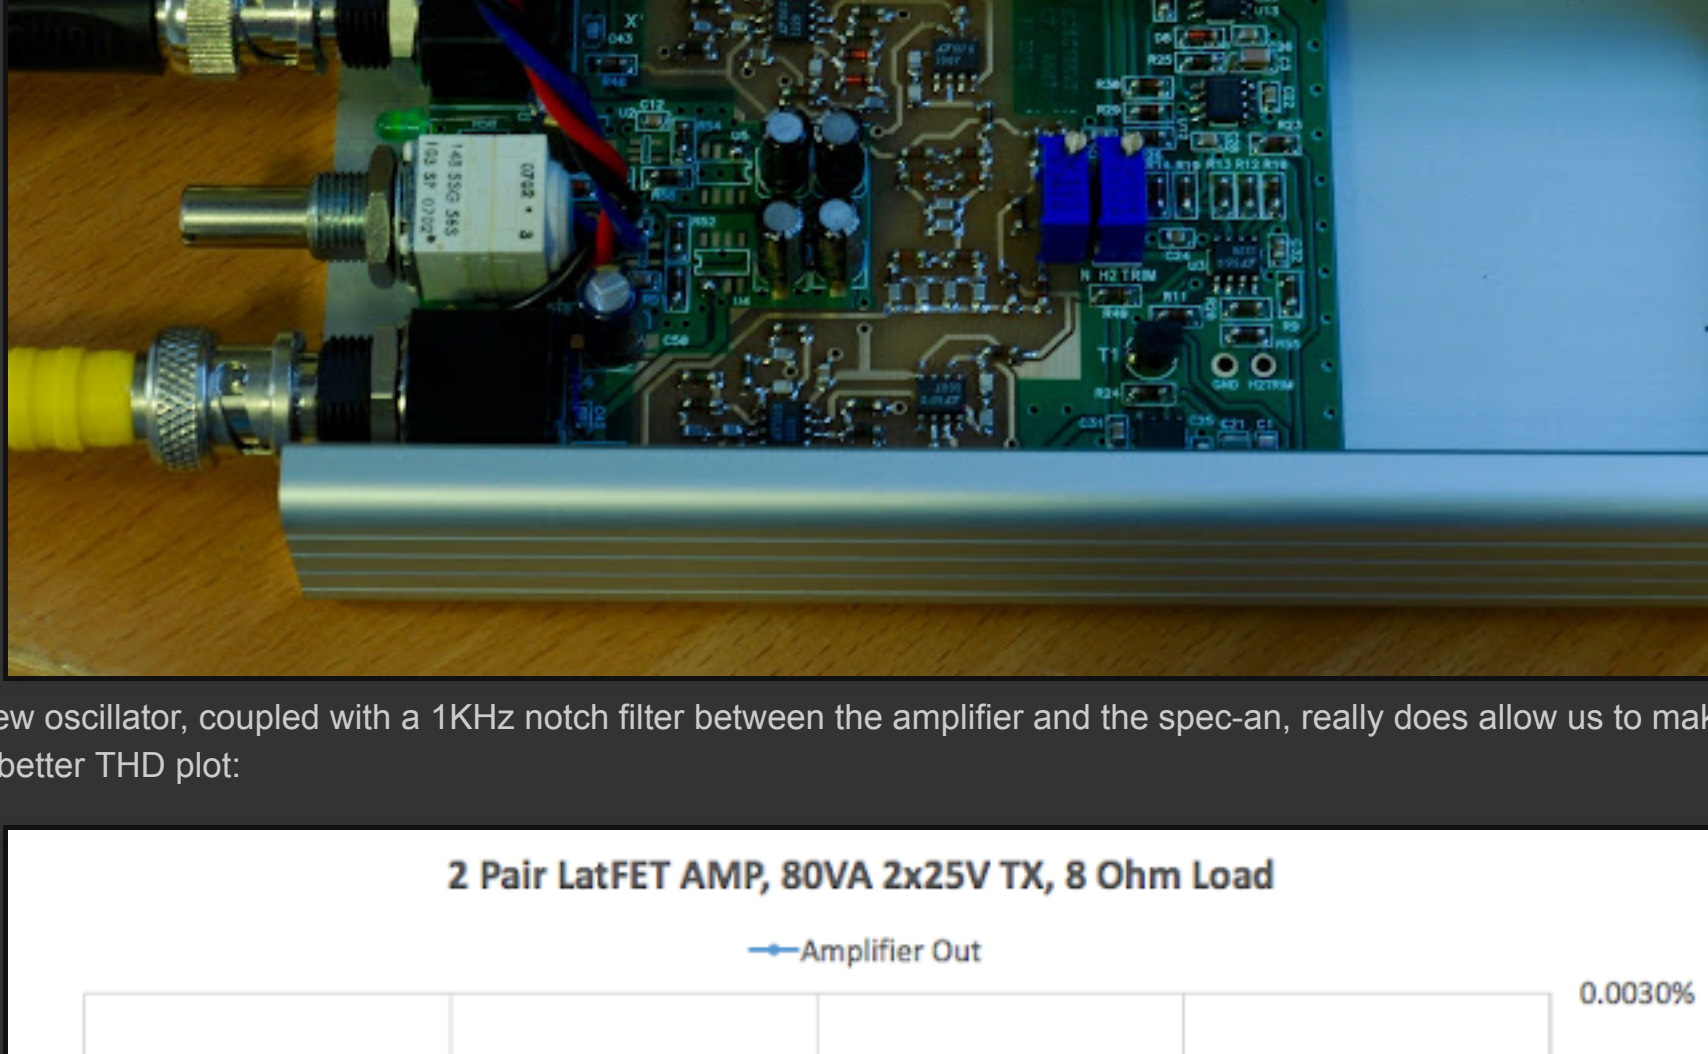
With the marker at 1KHz, we read off 139nV/sqrtHz. Input referred (this amp has a gain of 16), that's just 8.7nV/sqrtHz. To put it into context, the OP134, a pretty well regarded low-noise opamp, has an input referred voltage noise density of 8nV/sqrtHz.

Next we start working our way down the THD rabbit hole. I started by using my macBook air, as it has a reasonably clean DAC. I used Audacity to generate a 1KHz tone, then played that at all the max volume settings, both directly and via my amp output:

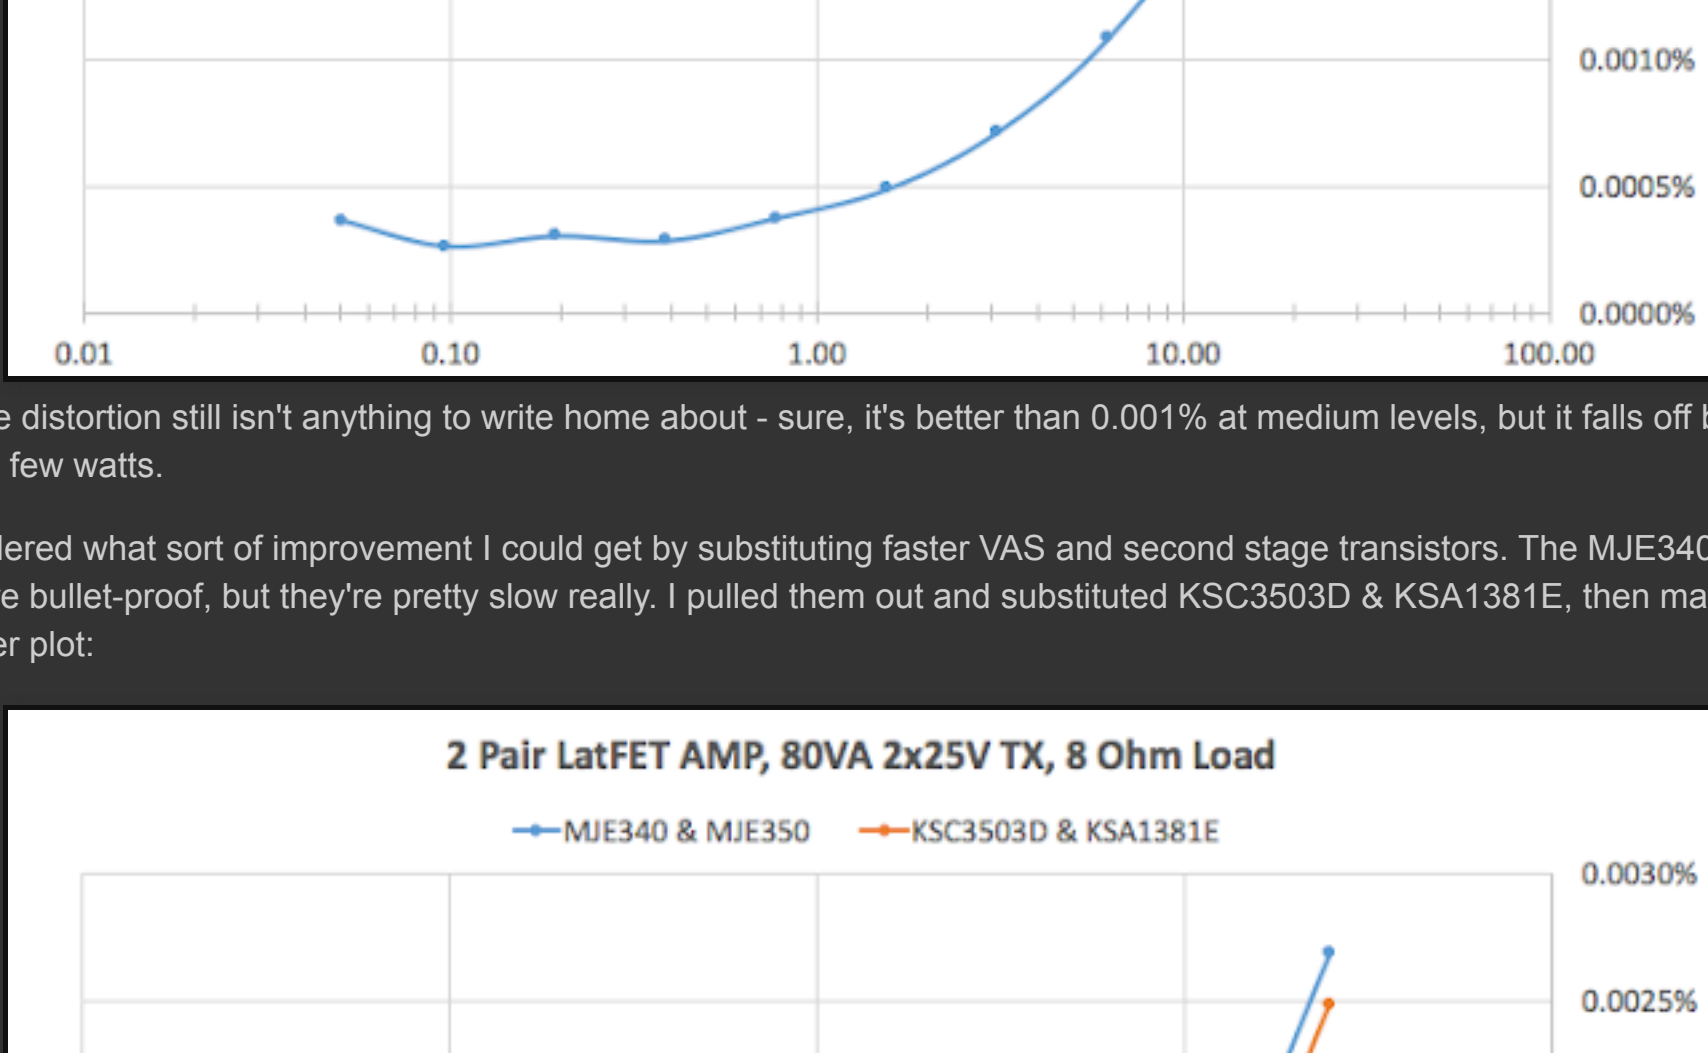


Okay, not bad. Note the amp appears better than the macBook around 100mW to 1W. This is simply because the amp provides a gain of 16, so what was buried in the noise of the spec-an without the amp, isn't with the amp. Things turn to custard after a few watts. Initially I passed that off on the wimpy 2x25V 80VA supply I was using for test, but we'll find later it's a much cooler problem.

To go deeper, I really do need a better source. So I spent some time building the 1KHz low-distortion oscillator. Here's a photo:

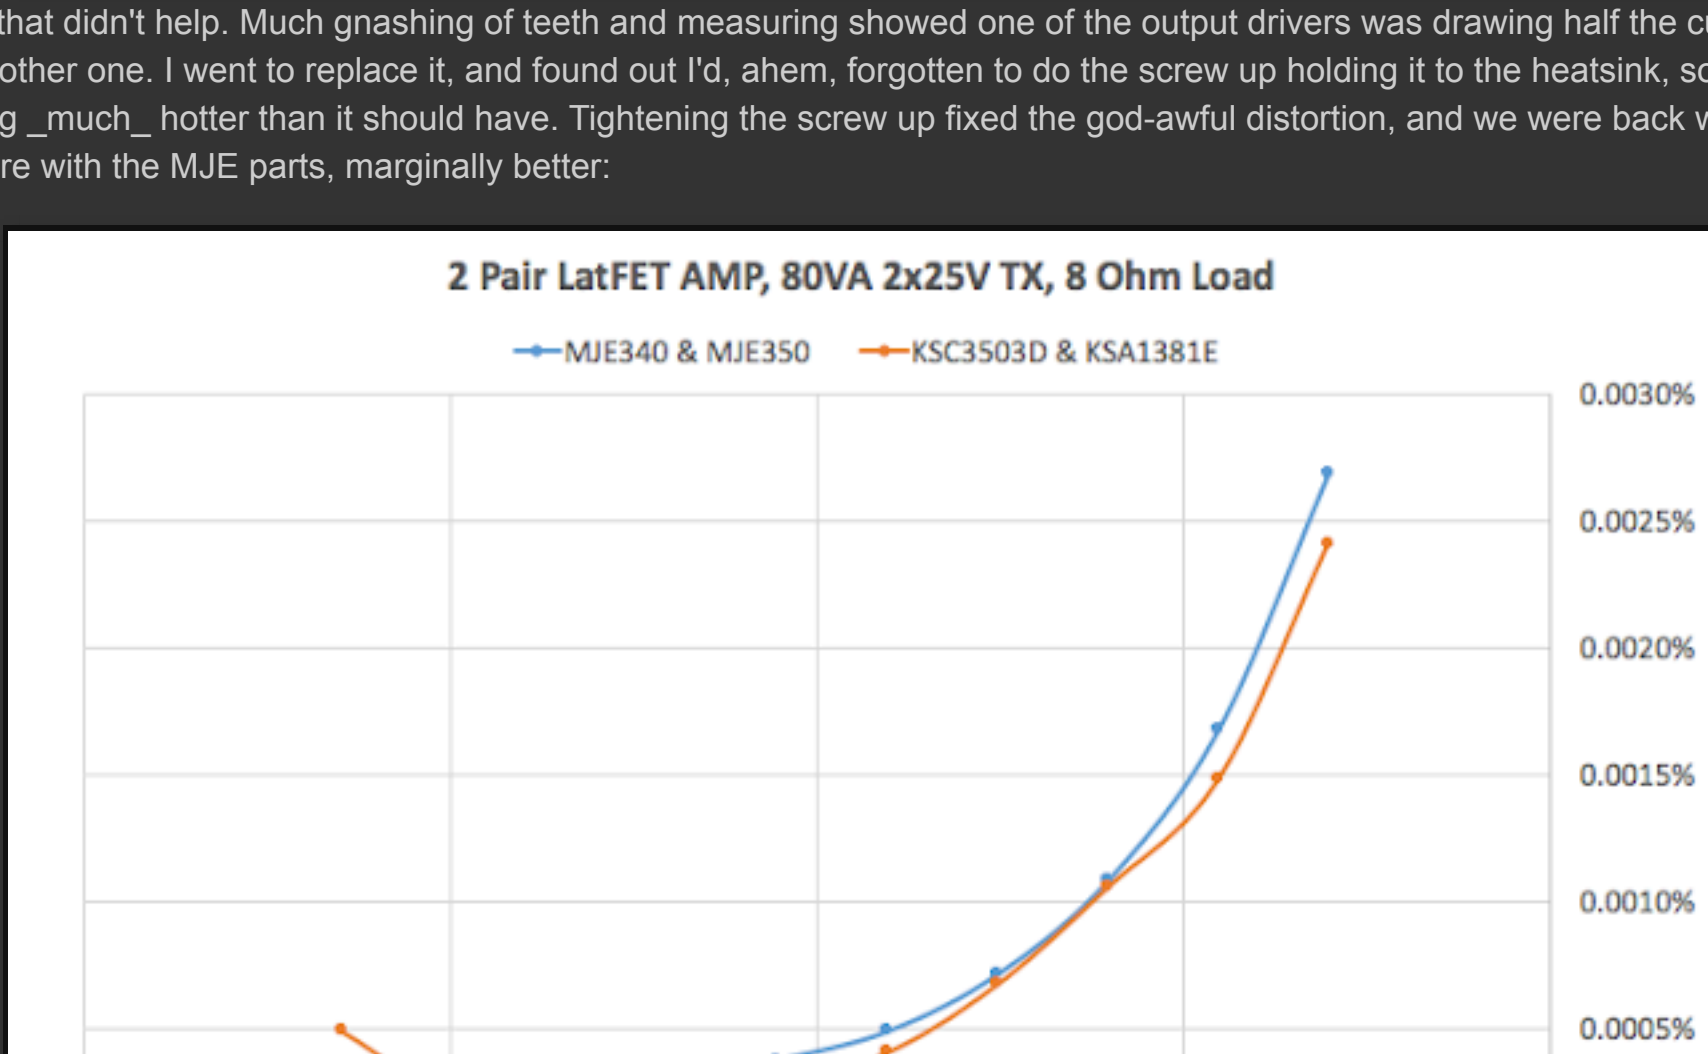


The new oscillator, coupled with a 1KHz notch filter between the amplifier and the spec-an, really does allow us to make a much better THD plot:




But the distortion still isn't anything to write home about - sure, it's better than 0.001% at medium levels, but it falls off badly after a few watts.

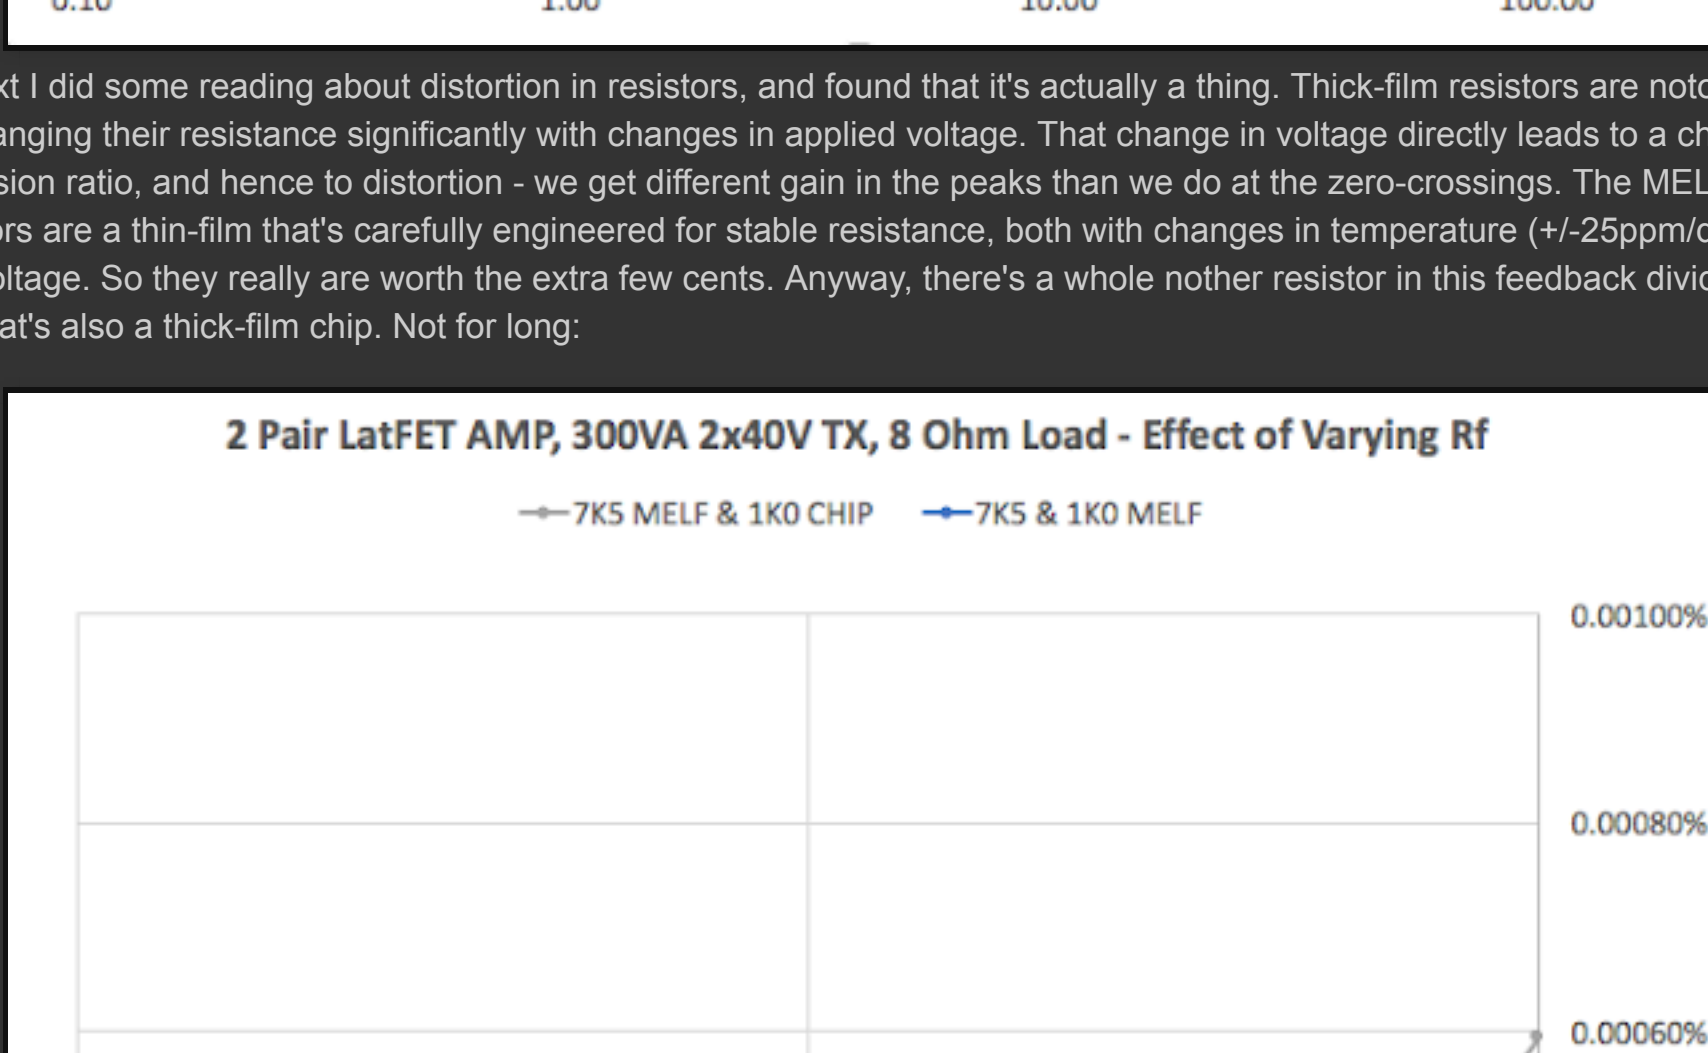
I wondered what sort of improvement I could get by substituting faster VAS and second stage transistors. The MJE340 and 360 are bullet-proof, but they're pretty slow really. I pulled them out and substituted KSC3030D & KSA1381E, then made another plot:



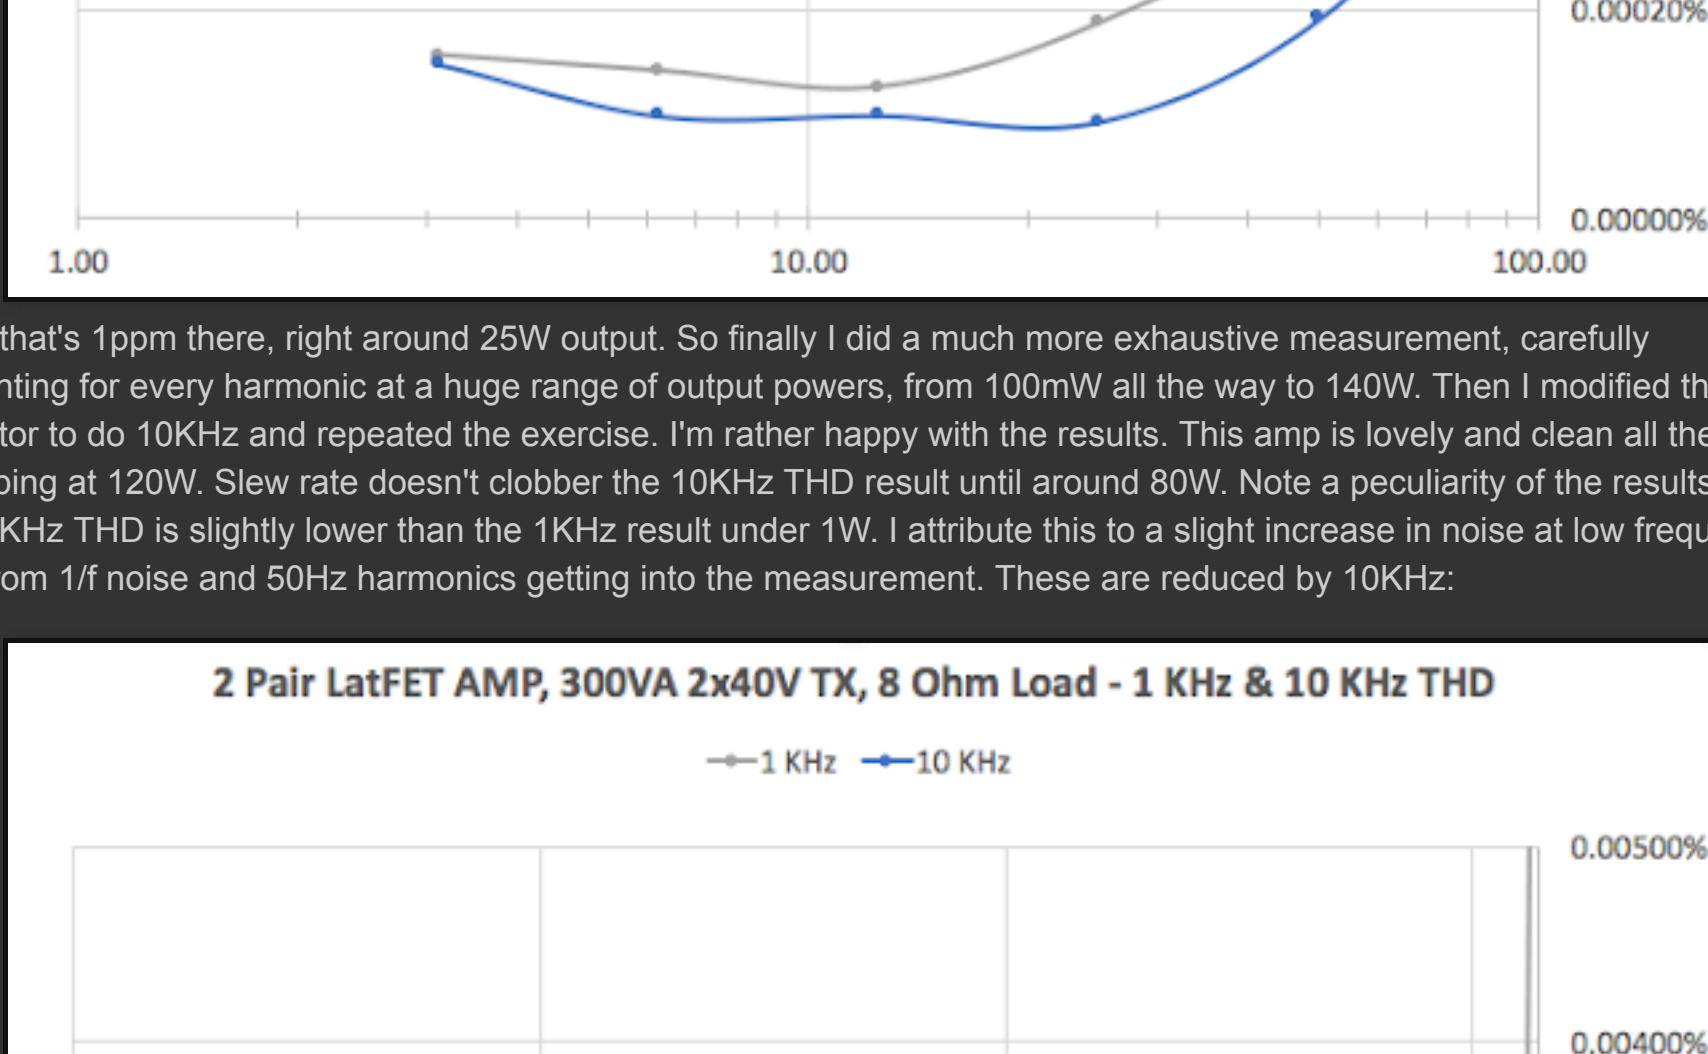
Wow, that didn't help. Much gnashing of teeth and measuring showed one of the output drivers was drawing half the current of the other one. I went to replace it, and found out I'd, ahem, forgotten to do the screw up hooking it to the heatsink, so it was running much hotter than it should have. Tightening the screw up fixed the god-awful distortion, and we were back where we were with the MJE parts, marginally better:



On the basis that the KSA and KSC parts are faster, and my sims showed considerably improved stability, I wondered if I could improve the THD a little by winding up the negative feedback. This has the effect of reducing the output level, but I can sort that with more gain in the preamp. So I substituted a 5K1 MELF resistor for the 15K thick-film 1200 that was setting the feedback ratio. Bam! Instant huge improvement in THD. Much better than the couple of dB that I should have gotten. So I tried a variety of other MELF resistors for Rf. All show dramatically improved THD over the 15K thick-film chip resistor, and a gradual increase in THD in line with reduction in feedback:



So next I did some reading about distortion in resistors, and found that it's actually a thing. Thick-film resistors are notorious for changing their resistance significantly with changes in applied voltage. That change in voltage directly leads to a change of division ratio, and hence to distortion - we get different gain in the peaks than we do at the zero-crossings. The MELF resistors are a thin-film that's carefully engineered for stable resistance, both with changes in temperature (+/-25ppm/deg) and voltage. So they really are worth the extra few cents. Anyway, there's a whole nother resistor in this feedback divider, and that's also a thick-film chip. Not for long:



Yeah, that's 1ppm there, right around 20W output. So finally I did a much more exhaustive measurement, carefully accounting for every harmonic at a huge range of output powers, from 100mW all the way to 140W. Then I modified the oscillator to 120W. Slow rate doesn't clobber the 10KHz THD result until around 80W. Note a peculiarity of the results where the 10KHz THD is slightly lower than the 1KHz result under 1W. I attribute this to a slight increase in noise at low frequency, both from 1/f noise and 50Hz harmonics getting into the measurement. These are reduced by 10KHz:



Posted by Suzy at 09:35

6 comments:

Raphael Pereira said...

Hello Suzy,

Finely today I found your web page again!

Let me make a simple question. Do you have some substitutes for the SST404, new part numbers pleased, and for the output mosfets.

I'm from Brazil, and it's very difficult to buy some parts.

The ones you used stay obsolete at this days.

Thanks for you attention!

11 April 2018 at 20:21

Suzy said...

CAlogic still make the SST404, and it's sold by Future Electronics.

http://www.futureelectronics.com/en/technologies/semiconductors/discretes/transistors/fets/Pages/8958308-SST404-LF.aspx?r=0

If you're doing your own build, the U401 is the same thing in a through-hole package, but with better match between the N and P-channel parts. To get roughly equivalent step response, I added a series R-C combination on the P-channel FET just to allow finer tweakability.

http://www.futureelectronics.com/en/technologies/semiconductors/discretes/transistors/fets/Pages/4318699-U401.aspx?r=0

12 April 2018 at 08:02

Suzy said...

I've been using Exicon ECX10P20 and ECX10N20 output devices. They're a drop-in replacement for the 2SK1058 & 2SJ162. They're distributed by Profusion. http://www.profusionscnc.com/type/ariel-mosfet

12 April 2018 at 08:16

Synonymous said...

Hi Suzy,

Can you explain the usage and how you came to the values of C9 C10 C11 C12 and R's 41 and 44 on the gates of the laterals please. Why are they different for both the P and N. In other works, I've seen either rod ellott use just a capacitor on the 1058, or both using a 15pf.

Thanks for any insight.

S

18 July 2020 at 02:02

Suzy said...

Hi S,

The rationale behind adding a capacitor to the N-channel MOSFET is simply to roughly balance the gate capacitance of the N and P-channel parts. To get roughly equivalent step response, I added a series R-C combination on the P-channel FET just to allow finer tweakability.

The values are not at all critical, and indeed as you've noted other amps omit them or use simpler schemes. I find that as the OLG gets up you need to pay more attention to small stuff like this to keep the amplifier from oscillating.

18 July 2020 at 02:12

Tim said...

Hi Suzy,

In the past I have seen and tested several power transistors 2N3055 and mosfets used in both switch mode power supplies and audio amplifiers in the past that had failed and found that there were replacement components in Australia which did not even come close to the manufacturers specifications, this along with cost of hobby electronics frustration of finding hard to get components and cheap imported items available seemed to produce a decline in hobby electronics in Australia over many years and is still happening

29 August 2020 at 20:13

Post a Comment

Newer Post

Subscribe to: [Post Comments \(Atom\)](#)

Home

Older Post

Simple theme. Powered by [Blogger](#).